

Node.js FS Module: Comprehensive Notes for Revision and Technical Interviews

This Markdown file provides detailed, technical yet simple notes on the Node.js `fs` (File System) module. Perfect for quick revision before interviews and hands-on practice. Every concept includes:

- **What** it is
 - **Why** we use it
 - **Real-world technical examples**
 - **Practice code snippets** (copy-paste and run in Node.js)
-

1. Introduction to the FS Module

What is the FS Module?

The `fs` module allows Node.js programs to interact with the operating system's file system. It supports full CRUD operations (Create, Read, Update, Delete) on files and directories.

Why Use the FS Module?

- Essential for server-side tasks: reading config files, logging errors, handling uploads, serving static files, etc.
- Gives direct, low-level control over the file system.
- Offers both **synchronous** (blocking) and **asynchronous** (non-blocking) methods to suit different use cases.

How to Import

```
// CommonJS
const fs = require('node:fs');

// ES Modules
import fs from 'node:fs';
```

Interview Tip: The `node:` prefix ensures you're loading the built-in module, not a third-party package with the same name.

2. Synchronous Operations (Blocking)

What & Why?

Synchronous methods end with `Sync` and **block** the event loop until the operation finishes.

Use them in:

- Simple scripts
- CLI tools
- Startup tasks where order matters and performance isn't critical

Avoid in production servers — blocking kills concurrency.

Key Methods

1. Create/Overwrite File — `writeFileSync`

```
fs.writeFileSync('./data.txt', 'Hello Sync!');  
console.log('File created/overwritten');
```

2. Read File — `readFileSync`

Returns Buffer by default. Use `'utf-8'` for string.

```
const content = fs.readFileSync('./data.txt', 'utf-8');  
console.log(content); // Hello Sync!
```

Why utf-8? Backward compatible with ASCII + variable-length encoding (efficient for most text).

3. Append to File — `appendFileSync`

Adds data at the end. Creates file if missing.

```
fs.appendFileSync('./data.txt', '\nNew line appended');
```

4. Delete File — `unlinkSync`

```
try {  
  fs.unlinkSync('./data.txt');  
  console.log('File deleted');  
} catch (err) {  
  if (err.code === 'ENOENT') console.log('File not found');  
}
```

5. Create Directory — `mkdirSync`

```
fs.mkdirSync('./logs/2025/12', { recursive: true }); // Creates nested folders
```

6. Remove Directory — `rmSync` (preferred) or `rmdirSync` (older)

```
fs.rmdirSync('./logs', { recursive: true, force: true });
```

7. Rename / Move File or Folder — `renameSync`

```
fs.renameSync('./old.txt', './new-name.txt'); // Rename
fs.renameSync('./file.txt', '../archive/file.txt'); // Move
```

8. Copy File — `copyFileSync`

```
fs.copyFileSync('./source.txt', './backup.txt');
```

3. Asynchronous Operations — Callbacks

What & Why?

Non-blocking. Use callbacks (error-first pattern).

Ideal for servers — keeps event loop free.

Key Methods

Read File

```
fs.readFile('./data.txt', 'utf-8', (err, data) => {
  if (err) return console.log(err);
  console.log(data);
});
console.log('This runs immediately!');
```

Write File

```
fs.writeFile('./async.txt', 'Async content', (err) => {
  if (err) console.log(err);
  console.log('File written');
});
```

Append File

```
fs.appendFile('./async.txt', '\nMore data', (err) => {
  if (err) console.log(err);
  console.log('Appended');
});
```

Ensuring Order (Avoid Callback Hell)

```
fs.writeFile('./order.txt', 'Step 1', (err) => {
  if (err) throw err;
  fs.appendFile('./order.txt', '\nStep 2', (err) => {
    if (err) throw err;
    fs.appendFile('./order.txt', '\nStep 3', () => {
      console.log('All steps done in order');
    });
  });
});
```

Interview Question:

Q: What is callback hell?

A: Deeply nested callbacks → hard to read/maintain. Solution: Promises or async/await.

4. Asynchronous Operations — Promises (`fs/promises`)

What & Why?

Cleaner than callbacks. Chain with `.then()` / `.catch()`.

```
const fsP = require('node:fs/promises');

// Write
fsP.writeFile('./promise.txt', 'Promise data')
  .then(() => console.log('Written'))
  .catch(err => console.log(err));

// Read
fsP.readFile('./promise.txt', 'utf-8')
  .then(data => console.log(data))
  .catch(err => console.log(err));

// Append
fsP.appendFile('./promise.txt', '\nAppended via promise');
```

5. Asynchronous Operations — Async/Await (Best for Modern Code)

What & Why?

Most readable. Looks synchronous but is non-blocking.

```
const fsP = require('node:fs/promises');

async function fileOps() {
  try {
    await fsP.writeFile('./await.txt', 'Start');
    await fsP.appendFile('./await.txt', '\nMiddle');
    await fsP.appendFile('./await.txt', '\nEnd');
    const content = await fsP.readFile('./await.txt', 'utf-8');
    console.log(content);
  } catch (err) {
    console.log('Error:', err);
  }
}

fileOps();
```

Interview Tip: Async functions return a Promise. `await` pauses execution inside the function but doesn't block the event loop.

6. Buffer in Node.js

What is Buffer?

Global class for handling raw binary data (fixed-size array of bytes).

Why Needed?

Files are binary. Text encoding (utf-8) converts to string, but sometimes you need raw bytes (images, videos, encryption).

```
const buf = Buffer.from('NodeJS');
console.log(buf);           // <Buffer 4e 6f 64 65 4a 53>
console.log(buf.toString()); // NodeJS

// Allocate empty buffer
const empty = Buffer.alloc(10);
empty.write('Hi');
console.log(empty.toString()); // Hi
```

Interview Question:

Q: `Buffer.alloc()` vs `Buffer.from()`?

A: `alloc(size)` creates zero-filled buffer. `from(data)` creates from string/array/etc.

7. Streams — The Most Efficient Way for Large Files

What are Streams?

Process data in chunks instead of loading everything into memory.

Why Use Streams?

- Low memory usage (critical for GB-sized files)
- Faster start (can begin processing first chunk immediately)
- Backpressure handling

Types

- **Readable** — `createReadStream`
- **Writable** — `createWriteStream`
- **Duplex** — both read & write
- **Transform** — modify data

Basic Readable Stream

```
const readStream = fs.createReadStream('./big-file.txt', {
  encoding: 'utf-8',
  highWaterMark: 1024 // 1KB chunks
});

readStream.on('data', (chunk) => {
  console.log('Chunk received:', chunk.length);
});
```

Piping (Best Way to Copy)

```
const read = fs.createReadStream('./source.txt');
const write = fs.createWriteStream('./copy.txt');

read.pipe(write); // Automatically handles chunks & backpressure
write.on('finish', () => console.log('Copy complete'));
```

Transform Example (Uppercase Converter)

```
const read = fs.createReadStream('./input.txt', { encoding: 'utf-8' });
const write = fs.createWriteStream('./upper.txt');

read.on('data', (chunk) => {
  write.write(chunk.toUpperCase());
});
```

Interview Question:

Q: Why use streams instead of `readFileSync` for large files?

A: `readFileSync` loads entire file → high memory usage + slow start. Streams use constant memory and start processing immediately.

1. What is the Node.js `fs` module and why is it important?

Answer:

The `fs` (File System) module provides an API to interact with the file system (read, write, delete files/folders). It's crucial for server-side applications that need to handle logs, configs, uploads, static assets, or any persistent data operations.

2. How do you import the `fs` module in Node.js?

Answer:

```
// CommonJS
const fs = require('node:fs');

// ES Modules
import fs from 'node:fs';
```

Using `node:` prefix ensures it's the built-in module.

3. What is the difference between synchronous and asynchronous `fs` methods?

Answer:

- Synchronous (`...Sync`) methods block the event loop until completion (e.g., `readFileSync`).
 - Asynchronous methods are non-blocking and use callbacks/promises (e.g., `readFile`).
Use `async` in servers for concurrency; `sync` in scripts for simplicity.
-

4. Why should you avoid synchronous `fs` methods in production servers?

Answer:

Node.js is single-threaded. Synchronous operations block the event loop, preventing it from handling other requests, timers, or I/O. This reduces throughput and can make the server unresponsive.

5. Explain `readFileSync` vs `readFile`.

Answer:

- `readFileSync`: Blocks execution, returns data directly.

- `readFile`: Non-blocking, takes a callback (or returns Promise in `fs/promises`).

Example: `readFile` allows the server to serve other clients while reading a file.

6. What happens if you don't specify encoding in `readFileSync` or `readFile`?

Answer:

It returns a `Buffer` object (raw binary data) instead of a string. You need to call `.toString()` or specify `'utf-8'` to get readable text.

7. Why is UTF-8 the most commonly used encoding?

Answer:

- Backward compatible with ASCII (English characters use 1 byte).
 - Variable-length encoding (1–4 bytes per character) — efficient for most languages.
 - Default in Node.js for `toString()` on Buffers.
-

8. What is a Buffer in Node.js?

Answer:

A `Buffer` is a global class for handling raw binary data directly. It's like an array of bytes, used when dealing with files, networks, images, or streams where text encoding isn't appropriate.

9. Difference between `Buffer.from()` and `Buffer.alloc()`?

Answer:

- `Buffer.from('hello')`: Creates buffer from data (string, array, etc.).
 - `Buffer.alloc(10)`: Creates a zero-filled buffer of fixed size (safe, no leftover data).
Use `alloc` for security when initializing empty buffers.
-

10. How do you append data to a file synchronously and asynchronously?

Answer:

Sync: `fs.appendFileSync(path, data)`

Async: `fs.appendFile(path, data, callback)` or `fsP.appendFile()` with promises.

11. How do you create nested directories in one command?

Answer:

Use `{ recursive: true }`:

```
fs.mkdirSync('a/b/c', { recursive: true });
```

12. How do you delete a non-empty directory?

Answer:

Use `fs.rmSync(path, { recursive: true, force: true })` (modern Node.js) or older `fs.rmdirSync` with recursive.

13. What is callback hell? How does it relate to `fs` operations?

Answer:

Callback hell occurs when multiple async `fs` operations are nested for sequential execution:

```
fs.readFile(..., () => {
  fs.writeFile(..., () => {
    // deeply nested
  });
});
```

Solution: Use Promises or `async/await`.

14. How do you use `fs/promises` for cleaner async code?

Answer:

```
const fs = require('node:fs/promises');
await fs.writeFile('file.txt', 'data');
await fs.appendFile('file.txt', 'more');
const data = await fs.readFile('file.txt', 'utf-8');
```

15. What are streams in Node.js? Why are they important?

Answer:

Streams allow processing data in chunks instead of loading everything into memory. Critical for large files to avoid memory exhaustion and enable faster processing.

16. Name the four types of streams in Node.js.

Answer:

- Readable (e.g., file input)
- Writable (e.g., file output)
- Duplex (read + write, e.g., sockets)

- Transform (modify data while streaming)
-

17. How do you create a readable and writable stream?

Answer:

```
const readStream = fs.createReadStream('big.txt');
const writeStream = fs.createWriteStream('copy.txt');
```

18. What is piping? Give an example.

Answer:

Piping connects a readable stream to a writable stream automatically:

```
fs.createReadStream('source.txt').pipe(fs.createWriteStream('dest.txt'));
```

Best way to copy large files efficiently.

19. What is **highWaterMark** in streams?

Answer:

It controls the maximum chunk size (in bytes) a readable stream buffers. Default is 64KB. Lower it for slower consumers to manage backpressure.

20. Why are streams better than **readFile** for large files?

Answer:

readFile loads entire file into memory → high RAM usage + delay.

Streams process in chunks → constant memory, immediate start, better performance.

21. How do you handle errors in asynchronous **fs** operations?

Answer:

Always check the first **err** parameter in callbacks:

```
fs.readFile('file.txt', (err, data) => {
  if (err) return console.error(err);
  // use data
});
```

With promises: use try/catch or **.catch()**.

22. What does `ENOENT` error mean?

Answer:

"Error NO ENTry" — file or directory does not exist. Common when reading/deleting non-existent paths.

23. How do you copy a file using `fs`?

Answer:

Best way: `fs.copyFileSync(src, dest)` (sync) or use streams with `.pipe()` (async).

24. How do you rename and move a file in one operation?

Answer:

`fs.renameSync(oldPath, newPath)` works for both renaming and moving (even across directories).

25. What is backpressure in streams?

Answer:

When a readable stream produces data faster than the writable stream can consume it. Node.js streams handle it automatically with buffering and pausing.

26. Can you transform data while streaming? Example?

Answer:

Yes, using `.on('data')` and modifying chunks:

```
readStream.on('data', chunk => {
  writeStream.write(chunk.toString().toUpperCase());
});
```

Or use built-in `Transform` streams.

27. When would you use `fs.watch()`?

Answer:

To monitor file/directory changes (e.g., auto-reload config in dev, file watchers). Note: Not 100% reliable on all OS.

28. What is the best practice for file operations in modern Node.js?

Answer:

Use `fs/promises` with `async/await` for clean, readable, non-blocking code. Use streams for large files.

29. Explain how to read a large JSON file efficiently.

Answer:

Use streams + JSON parser (like `JSONStream` or manual chunk parsing). Avoid `readFileSync` + `JSON.parse()` — it loads everything into memory.

30. If a file operation fails halfway, how do you ensure data consistency?

Answer:

- Write to a temporary file, then `fs.rename()` (atomic on same filesystem).
 - Use transactions (if using databases).
 - Or implement rollback logic manually.
-

Tip for Interview: Always mention **non-blocking I/O, event loop impact, and memory efficiency** when discussing `fs` operations. These show deep understanding of Node.js architecture.

Key Interview Tips Summary

Topic	Key Point to Remember
Sync vs Async	Sync blocks event loop → bad for servers
Callback Hell	Nested callbacks → use Promises or async/await
UTF-8	Most popular encoding: backward compatible + efficient
Buffer	For raw binary data (images, crypto, etc.)
Streams	Best for large files: low memory, fast start, backpressure handling
Best Practice	Use <code>fs/promises</code> + async/await in modern Node.js